Demo Project: Deploy Web Application on EC2 Instance (manually)

This guide demonstrates how to:

1. Create and configure an EC2 instance on AWS.
2. Connect to the EC2 instance via SSH.
3. Install Docker on the remote EC2 instance.
4. Deploy a Docker container from a private DockerHub repository.
5. Make the application accessible publicly.

---

1. Create and Configure an EC2 Instance on AWS
2. Install Docker on the EC2 Instance
3. Deploy the Docker Image from a Private Docker Repository
4. Make the Application accessible from the Browser
    Troubleshooting

---

1. Create and Configure an EC2 Instance on AWS

1. **Log in to the AWS Management Console**:
    - Visit AWS Console.
2. **Launch an EC2 Instance**:
    - Navigate to **EC2 Dashboard** → **Launch Instances**.
    - Configure the following:
        - **Name**: `my-instance`.
        - **Tags**: Add tags for better resource management:

        - **Key:** `Type`
        - **Value:** `web-server-with -docker`
    - **AMI (Amazon Machine Image)**: Use **Amazon Linux 2023 AMI**.
    - **Instance Type**: Select **t2.micro** (Free Tier eligible).
    - **Key Pair**: Create or select an existing **SSH key pair** for secure access. For this project create a new Key Pair as `docker-server.pem`.
    - VPC: Choose VPC if you have a specific VPC that EC2 needs to be deployed, else leave in default VPC.
    - Subnet: Choose the subnet if you have a specific VPC that EC2 needs to be deployed, else leave as “No preference”
    - **Network Settings**:
        - **VPC**: Select your specific VPC or leave it as **Default VPC**.
        - **Subnet**: Choose a specific subnet if applicable or leave it as **No Preference**.
        - Auto-assign a **public IP**.
        - Add a **Security Group** with:
            - Add **Inbound Rules** for:
                - **SSH (Port 22)**: Your IP only (e.g., `203.0.113.0/32`).
    - **Storage**: Allocate at least **8GB** SSD (GP3).
- Click **Launch Instance**.

3. Move the Key Pair to .ssh: `mv Downloads/docker-server.pem ~/.ssh/`

   Confirm it was moved successfully: `ls -l .ssh`
4. Update Key Permissions: `chmod 400 .ssh/docker-server.pem`
5. **Access the EC2 Instance via SSH**:
    - Copy the **public IP address** of the instance.
    - Use the following command to connect:

      `ssh -i <path-to-your-private-key.pem> ec2-user@<public-ip-address>`

Example: `ssh -i .ssh/docker-server.pem ec2-user@52.71.72.116`

---

2. Install Docker on the EC2 Instance

1. **Update the Package List**: `sudo yum update -y`
2. **Install Docker**: `sudo yum install docker -y`
3. **Verify Docker Installation**: `docker --version`
4. **Start Docker**: `sudo service docker start`
5. Check docker is running by checking it's process: `ps aux | grep docker`
6. **Add the EC2 User to the Docker Group**:
   `sudo usermod -aG docker ec2-user`
   - **Why This Step is Necessary**:
     - The Docker daemon runs as the **root user**.
     - Regular users (like `ec2-user`) cannot execute Docker commands unless they are part of the **Docker group**.
     - Adding the EC2 user (`ec2-user`) to the Docker group with the `usermod` command grants it the necessary permissions to run Docker commands **without using sudo**.
   - **Apply Changes**: Log out and log back in for the changes to take effect:
     ```
     exit
     ssh -i <path-to-your-private-key.pem> ec2-user@<public-ip-address>
     ```

---

3. Deploy the Docker Image from a Private Docker Repository

1. **Log in to Your Private Docker Repository**:
   - Replace `<username>` and `<password>` with your DockerHub credentials or token:
     `docker login -u <username> -p <password>`

2. **Pull the Docker Image**:
    - Replace `<image-name>` with your private repository's image name

      `docker pull <username>/<image-name>:<tag>`

      Example: `docker pull eduardobautistamaciel/demo-app:1.0`
3. **Run the Docker Container**:
    - Replace `<port>` with the port your application uses:

      `docker run -d -p <port>:<container-port> <username>/<image-name>:<tag>`

      Example: `docker run -d -p 3000:3080 eduardobautistamaciel/demo-app:1.0`
4. Verify the container is running: `docker ps`

---

4. Make the Application accessible from the Browser

1. **Update the Security Group**:

- Go to **AWS EC2 → Security Groups → Edit Inbound Rules**.
- Add a rule:
    - **Type**: Custom TCP Rule.
    - **Port**: 3000.
    - **Source**: Anywhere (`0.0.0.0/0`).

2. **Verify the Application**:
    - Open your browser and navigate to:

      `http://<public-ip-address>:<port>`
    - Example: `http://52.71.72.116:3000`

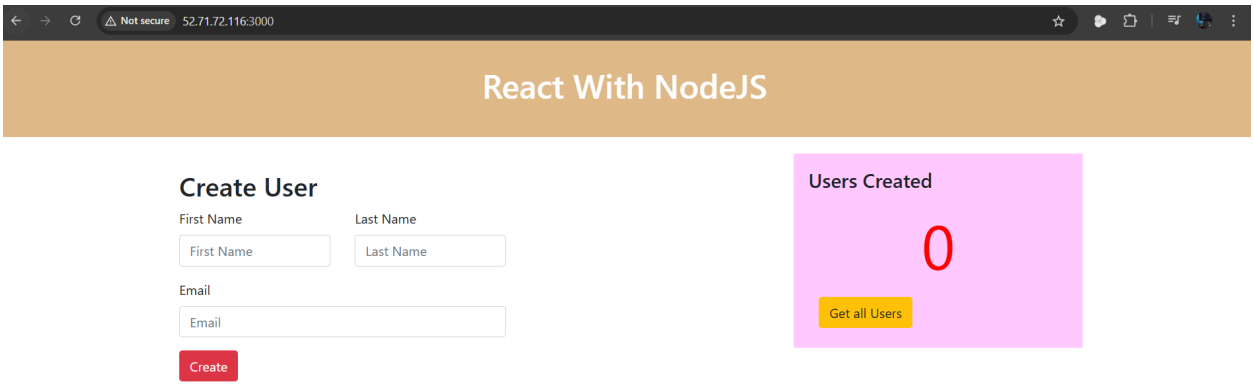

Troubleshooting

1. **Cannot Connect to EC2 Instance**:
   - Verify the **Security Group rules** allow SSH access from your IP.
   - Ensure your key pair matches the instance's key pair.
2. **Docker Permission Denied**:
   - **Cause**: Docker commands require root privileges.
   - **Solution**: Add the EC2 user to the Docker group (`usermod -aG docker ec2-user`), log out, and log back in to apply the group change.
3. **Application Not Accessible in Browser**:
   - Ensure the **Security Group** allows traffic on port `3000`.
   - Verify the container is running (`docker ps`).
   - Check the container logs: `docker logs <container-id>`
4. **Cannot Pull Docker Image**:
   - Verify your credentials or token for the private Docker repository.
   - Ensure the image name and tag are correct.