

# Demo Project: Deploy Microservices application in Kubernetes with Production & Security Best Practices

This guide demonstrates deploying an online shop application with **10 microservices** and **Redis** into a Kubernetes cluster hosted on Linode. It also includes production-grade configurations and security best practices.

---

## Step 1: Gather Developer-Provided Information

Example Architecture:

Step 2: Create Deployment & Service Configurations

Step 3: Set Up the Kubernetes Cluster

Step 4: Configure Access to the Cluster

Step 5: Deploy Microservices

Step 6: Verify Application Access

Step 7: Make Sure to Consider Security Best Practices

1. Container Image Versioning
2. Probes for Application Health
3. Resource Requests and Limits
4. High Availability
5. Organizational Practices
6. Security Measures
7. Networking Security

Step 8: Clean-Up

GitHub Security Tip:

---

## Step 1: Gather Developer-Provided Information

Before starting, ensure you have the following details from the developer:

## **1. Microservices Information:**

- Names of all microservices.
- Docker image names and versions.

## **2. Configuration Details:**

- Ports exposed by each service.
- Required environment variables.

## **3. Connections:**

- Dependencies between microservices.
- External services or databases.

## **4. Resources:**

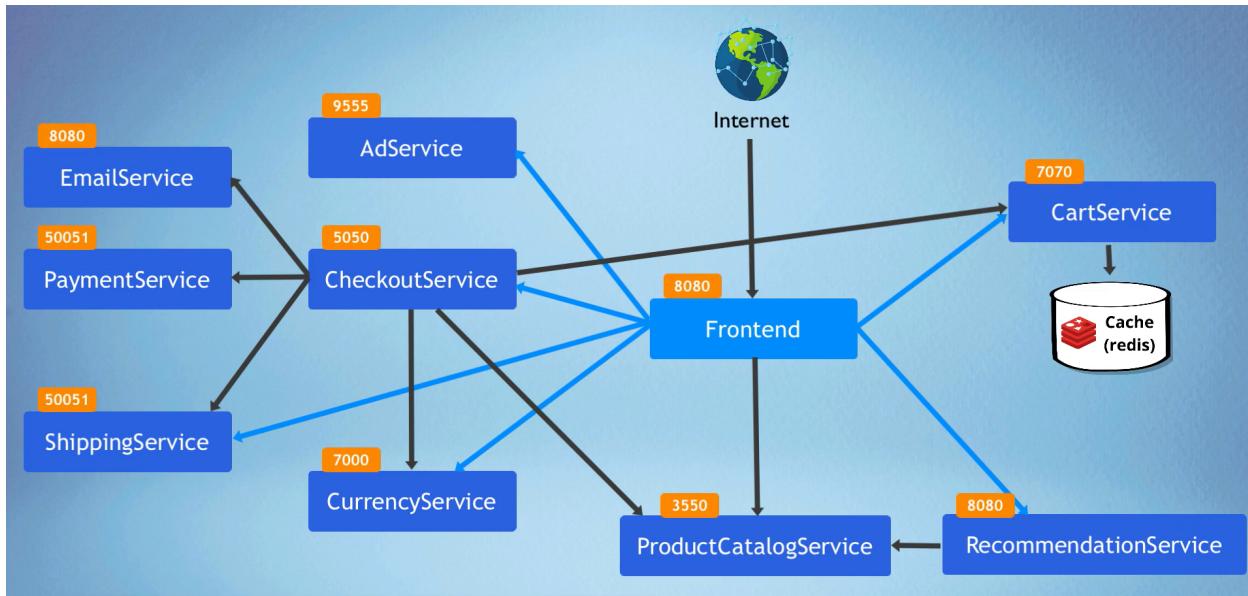
- CPU and memory requirements.
- Limits for each container.

## **5. Accessibility:**

- Identify which services will be exposed outside the cluster.

## **Example Architecture:**

The online shop application consists of 10 microservices and Redis for storing shopping cart data. All services are in the same namespace (`microservices`) since one team manages them.



## Step 2: Create Deployment & Service Configurations

### 1. Prepare the `config.yaml`:

- The file should include **Deployment** and **Service** manifests for each microservice.
- For **each microservice**, define:
  - **Environment Variables**.
  - **Ports** for communication.
  - **Resource Requests and Limits** for containers.
  - **Liveness and Readiness Probes**.

### 2. Example Configuration for a Single Microservice:

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: emailservice
spec:
  selector:
    matchLabels:
      app: emailservice
  template:
    metadata:
      labels:
        app: emailservice
  spec:
    containers:
      - name: service
        image: gcr.io/google-samples/microservices-demo/emailservice:v0.8.0
        ports:
          - containerPort: 8080
        env:
          - name: PORT
            value: "8080"
        livenessProbe:
          grpc:
            port: 8080
            periodSeconds: 5
        readinessProbe:
          grpc:
            port: 8080
            periodSeconds: 5
        resources:
          requests:
            cpu: 100m
            memory: 64Mi
          limits:
            cpu: 200m
            memory: 128Mi
    ---
apiVersion: v1
kind: Service
metadata:
  name: emailservice
  type: ClusterIP
  selector:
    app: emailservice
  ports:
    - protocol: TCP
      port: 5000
      targetPort: 8080

```

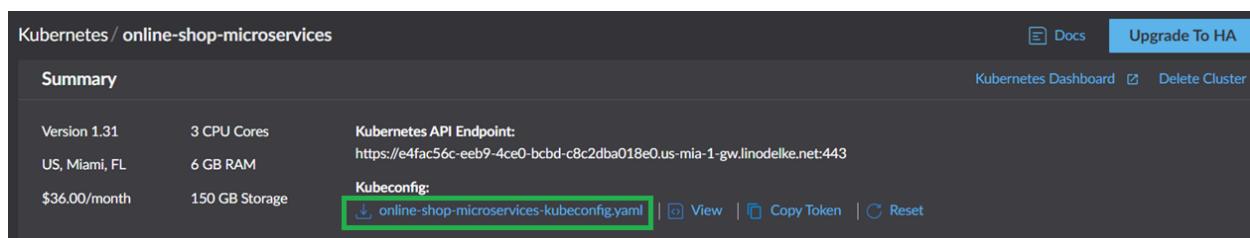
# Step 3: Set Up the Kubernetes Cluster

## 1. Create a Kubernetes Cluster in Linode:

- Navigate to **Create → Kubernetes** in the Linode dashboard.
- **Cluster Details:**
  - **Cluster Label:** `online-shop-microservices`
  - **Region:** Select the closest region.
  - **Kubernetes Version:** Use the latest version.
- **HA Control Plane:** No
- **Node Pools:** Add at least **3 nodes** (2GB Shared CPU each).
- Click **Create Cluster**.

## 2. Download the Kubeconfig File:

- Save the `online-shop-microservices-kubeconfig.yaml` file after the cluster is created. This file is required for `kubectl` to interact with the cluster.



# Step 4: Configure Access to the Cluster

## 1. Secure Kubeconfig File:

- Place the `online-shop-microservices-kubeconfig.yaml` in your working directory.
- Restrict access: `chmod 400 online-shop-microservices-kubeconfig.yaml`

## 2. Export the file:

```
export KUBECONFIG=online-shop-microservices-kubeconfig.yaml
```

### 3. Verify the Cluster:

- Confirm that the nodes are running: `kubectl get node`

## Step 5: Deploy Microservices

### 1. Create a Namespace for the Application: `kubectl create ns microservices`

### 2. Deploy the microservices:

- Apply the `config.yaml` manifest to the `microservices` namespace:

```
kubectl apply -f config.yaml -n microservices
```

### 3. Verify the Deployment:

- Check the status of pods: `kubectl get pod -n microservices`

Example:

NAME	READY	STATUS	RESTARTS	AGE
adservice-5c98bd5cb-qkvnn	1/1	Running	0	19h
cartservice-76c4d44889-s7wf8	1/1	Running	0	19h
checkoutservice-7cdbddccb8-t7bvk	1/1	Running	0	19h
currencyervice-9cc8864fd-qczsg	1/1	Running	0	19h
emailservice-7c85cfb958-lxbpk	1/1	Running	0	19h
frontend-6455d5bdc6-72tq5	1/1	Running	0	19h
frontend-6455d5bdc6-jvfnm	1/1	Running	0	19h
paymentservice-66f455f78b-tkdf5	1/1	Running	0	19h
productcatalogservice-6dc5664f4-9kzlp	1/1	Running	0	19h
recommendationservice-79d9b8d6d4-h5lrh	1/1	Running	0	19h
redis-cart-7d4fbca49bf-gqqt6	1/1	Running	0	19h
shippingservice-b468854fb-7r5p5	1/1	Running	0	19h

- Check services: `kubectl get svc -n microservices`

Example:

kubectl get svc -n microservices						
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	
adservice	ClusterIP	10.128.54.59	<none>	9555/TCP	19h	
cartservice	ClusterIP	10.128.54.97	<none>	7070/TCP	19h	
checkoutservice	ClusterIP	10.128.233.44	<none>	5050/TCP	19h	
currencyervice	ClusterIP	10.128.78.119	<none>	7000/TCP	19h	
emailservice	ClusterIP	10.128.36.25	<none>	5000/TCP	19h	
frontend	NodePort	10.128.241.15	<none>	80:30007/TCP	19h	
paymentservice	ClusterIP	10.128.33.250	<none>	50051/TCP	19h	
productcatalogservice	ClusterIP	10.128.133.62	<none>	3550/TCP	19h	
recommendationservice	ClusterIP	10.128.1.37	<none>	8080/TCP	19h	
redis-cart	ClusterIP	10.128.15.109	<none>	6379/TCP	19h	
shippingservice	ClusterIP	10.128.122.161	<none>	50051/TCP	19h	

## Step 6: Verify Application Access

### 1. Obtain an External IP Address:

- In the Linode dashboard, navigate to **NodeBalancers** and copy the external IP address of one of the node pools.

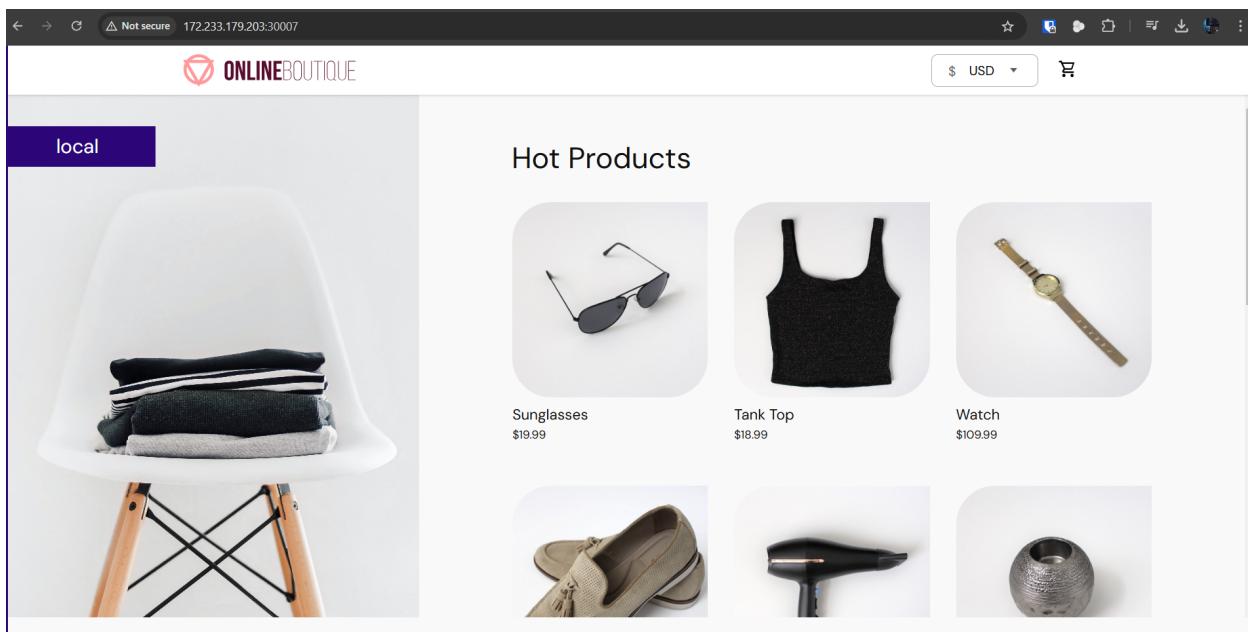
The screenshot shows the Kubernetes Dashboard for the `online-shop-microservices` cluster. In the **Summary** section, it displays cluster details: Version 1.31, 3 CPU Cores, 6 GB RAM, and 150 GB Storage. The Kubernetes API Endpoint is listed as `https://277ef1d3-089f-42b3-91f8-7d19b4e33658.us-mia-2.gw.linodelke.net:443`. The **Node Pools** section shows a pool named `Linode 2 GB` with three nodes. The IP addresses for these nodes are highlighted with a green box: `172.233.179.203`, `172.233.179.244`, and `172.233.179.98`.

Linode 2 GB		Autoscale Pool	Resize Pool	Recycle Pool Nodes	Delete Pool
Linode ^	Status ▾	IP Address ▾			
Ike323307-520771-00b788620000	Running	172.233.179.203 ⓘ			Recycle
Ike323307-520771-09a4ab380000	Running	172.233.179.244 ⓘ			Recycle
Ike323307-520771-38efbee50000	Running	172.233.179.98 ⓘ			Recycle

### 2. Access the Application in a Browser:

- Use the IP address and NodePort of the service exposed to the browser (e.g., `30007`):

`http://<node-pool-ip>:30007`



## Step 7: Make Sure to Consider Security Best Practices

### 1. Container Image Versioning

- **Pin (Tag) Versions** for each container image.
  - This ensures control and visibility of which versions of your applications are running inside the cluster.
  - Avoid using the `latest` tag, as it can lead to unexpected updates and instability.

### 2. Probes for Application Health

- **Liveness Probe:**

- Automatically restarts pods when the application crashes or experiences issues.
- It periodically pings the application endpoint (e.g., every 5 or 10 seconds) to confirm it's responding.

- Configured to run **after** the container has started while the application is running.
  - **Readiness Probe:**
    - Verifies that the application has fully started and is ready to receive traffic.
    - Ensures Kubernetes only routes requests to pods that are fully initialized.
- 

### 3. Resource Requests and Limits

- **Why?**
  - Avoid over-allocating or under-allocating cluster resources.
  - Prevent a single container from consuming all available resources.
- **Example Configuration:**

```
resources:
  requests:
    memory: "128Mi"
    cpu: "500m"
  limits:
    memory: "256Mi"
    cpu: "1"
```

**Note:** Ensure the requested values fit within the capacity of your cluster nodes. If the requested resources exceed the node size, the pod will not be scheduled.

---

### 4. High Availability

- **Replicas for Deployments:**
  - Use more than one replica to ensure availability. If one pod crashes, the application remains accessible.
  - Example: `replicas: 2`

- **Multiple Worker Nodes:**
    - Distribute replicas across multiple worker nodes to prevent a **Single Point of Failure**.
    - Reasons for server unavailability:
      - Server crashes.
      - Server reboots (updates or maintenance).
      - Hardware failures.
- 

## 5. Organizational Practices

- **Use Labels for Resources:**
    - Group resources (e.g., pods, services) using labels for better organization and reference in service components.
  - **Namespaces:**
    - Isolate resources using namespaces.
    - Simplifies management and enables role-based access control (RBAC) for different teams.
- 

## 6. Security Measures

### 1. Vulnerability-Free Images:

- Perform vulnerability and security scans on images manually or through your CI/CD pipeline.

### 2. No Root Access:

- Avoid running containers as root to reduce the risk of resource abuse or damage to the host.
- Official images usually avoid root access, but always verify third-party images.

### 3. Update Kubernetes:

- Regularly update Kubernetes to the latest version to address security vulnerabilities and bugs.
- 

## 7. Networking Security

- **Avoid NodePorts:**
  - Exposing a NodePort opens a port on all worker nodes, increasing the attack surface.
  - Instead, use **LoadBalancer** type services or an **Ingress Controller** for a secure single entry point.
- **Example Transition:**
  - Change frontend service type from **NodePort** to **LoadBalancer**:

### Before:

```
type: NodePort
nodePort: 30007
```

### After:

```
type: LoadBalancer
```

## Step 8: Clean-Up

### 1. Delete the Application:

- Remove all resources for the microservices:

```
kubectl delete -f config.yaml -n microservices
```

### 2. Testing LoadBalancer Transition:

- Apply changes to use **LoadBalancer** instead of **NodePort**, and update the deployment to use `replicas: 2` instead of `replicas: 1` for improved high availability:

- Update the frontend service configuration:

```
type: LoadBalancer
```

- Update the deployment configuration:

```
replicas: 2
```

- Reapply the updated configuration:

```
kubectl apply -f config.yaml -n microservices
```

### 3. Verify the Deployment:

Check the status of pods: `kubectl get pod`

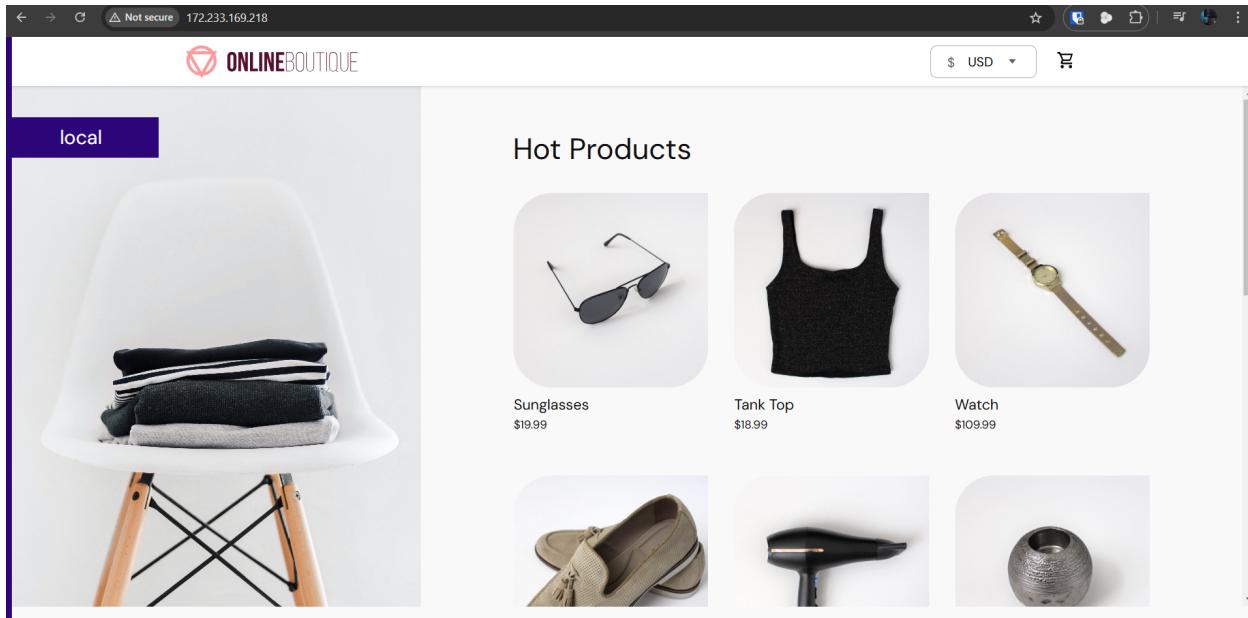
Example:

<code>kubectl get pod</code>	READY	STATUS	RESTARTS	AGE
adservice-54b74b55b7-ft6zg	1/1	Running	0	4m52s
adservice-54b74b55b7-xz8st	1/1	Running	0	4m53s
cartservice-dd4f7764f-wqpd7	1/1	Running	0	4m52s
cartservice-dd4f7764f-zj2l8	1/1	Running	0	4m51s
checkoutservice-767cc65db4-8w8cf	1/1	Running	0	4m54s
checkoutservice-767cc65db4-9bj4k	1/1	Running	0	4m54s
currencyervice-7dcb6cb45f-blmttd	1/1	Running	0	4m51s
currencyervice-7dcb6cb45f-n46qf	1/1	Running	0	4m51s
emailservice-6d488b67f8-5lt7z	1/1	Running	0	4m52s
emailservice-6d488b67f8-rhxrr	1/1	Running	0	4m53s
frontend-696858b68d-47q2w	1/1	Running	0	4m54s
frontend-696858b68d-nftjt	1/1	Running	0	4m54s
paymentservice-795d5d4bd8-7ggbx	1/1	Running	0	4m54s
paymentservice-795d5d4bd8-7m7zf	1/1	Running	0	4m53s
productcatalogservice-7cf54479b6-nlnc4	1/1	Running	0	4m52s
productcatalogservice-7cf54479b6-wxkr4	1/1	Running	0	4m53s
recommendationservice-76887dcc55-j8vvm	1/1	Running	0	4m54s
recommendationservice-76887dcc55-xx7ck	1/1	Running	0	4m54s
redis-cart-6dfd6dfd75-68pk7	1/1	Running	0	4m54s
shippingservice-d84b79bb-5ck2p	1/1	Running	0	4m54s
shippingservice-d84b79bb-kgdsv	1/1	Running	0	4m54s

### 4. Verify the application:

- Get the **LoadBalancer IP** from Linode NodeBalancer.

- Access the application in the browser using: `http://<LoadBalancer-IP>`



## 6. Final Clean-Up:

- Once testing is complete, delete all resources:

```
kubectl delete -f config.yaml -n microservices
```

## GitHub Security Tip:

- Exclude the `online-shop-microservices-kubeconfig.yaml` from version control by adding it to `.gitignore`

### Why?

- This file contains sensitive credentials and certificates that can compromise your cluster if exposed.