

Demo Project: Automate AWS Infrastructure

This project demonstrates how to automate the provisioning of AWS infrastructure using Terraform. The infrastructure components include a VPC, Subnet, Route Table, Internet Gateway, Security Group, and an EC2 instance running a Docker container. The Terraform script will also deploy a Docker container (e.g., running Nginx) on the EC2 instance automatically via `user_data`.

Step 1: Configure Terraform Variables

Step 2: Create VPC & Subnet

Step 3: Create Route Table and Internet Gateway

Step 4: Create Security Group

Step 5: Obtain the Latest Amazon Linux 2 AMI

Step 6: Create SSH Key Pair

Step 7: Create the EC2 Instance

Step 8: Deploy the Infrastructure

Step 9: Provisioners Branch (Feature: `feature/provisioners`)

Step 10: Clean Up

Step 1: Configure Terraform Variables

Create a file named `terraform.tfvars` with content similar to:

```
vpc_cidr_blocks    = "10.0.0.0/16"
subnet_cidr_block  = "10.0.10.0/24"
avail_zone         = "us-east-1a"
env_prefix         = "dev"
my_ip              = "167.57.113.78/32"
instance_type      = "t2.micro"
public_key_location = "/home/eb/.ssh/id_rsa.pub"
private_key_location = "/home/eb/.ssh/id_rsa"
```

Step 2: Create VPC & Subnet

Your **main.tf** includes the following:

- **VPC Resource:**

```
resource "aws_vpc" "myapp-vpc" {
  cidr_block = var.vpc_cidr_blocks
  tags = {
    Name = "${var.env_prefix}-vpc"
  }
}
```

- **Subnet Resource:**

```
hcl
CopyEdit
resource "aws_subnet" "myapp-subnet-1" {
  vpc_id      = aws_vpc.myapp-vpc.id
  cidr_block  = var.subnet_cidr_block
  availability_zone = var.avail_zone
  tags = {
    Name = "${var.env_prefix}-subnet-1"
  }
}
```

Step 3: Create Route Table and Internet Gateway

- **Internet Gateway:**

```
resource "aws_internet_gateway" "myapp-igw" {
  vpc_id = aws_vpc.myapp-vpc.id
  tags = {
```

```
Name = "${var.env_prefix}-igw"
}
}
```

- **Default Route Table:**

```
resource "aws_default_route_table" "main-rtb" {
  default_route_table_id = aws_vpc.myapp-vpc.default_route_table_id

  route {
    cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.myapp-igw.id
  }
  tags = {
    Name = "${var.env_prefix}-main-rtb"
  }
}
```

Step 4: Create Security Group

- **Default Security Group:**

```
resource "aws_default_security_group" "default-sg" {
  vpc_id = aws_vpc.myapp-vpc.id

  ingress {
    from_port = 22
    to_port   = 22
    protocol  = "TCP"
    cidr_blocks = [var.my_ip]
  }

  ingress {
    from_port = 8080
  }
}
```

```

    to_port    = 8080
    protocol   = "TCP"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port    = 0
    to_port      = 0
    protocol     = "-1"
    cidr_blocks  = ["0.0.0.0/0"]
    prefix_list_ids = []
  }

  tags = {
    Name = "${var.env_prefix}-default-sg"
  }
}

```

Step 5: Obtain the Latest Amazon Linux 2 AMI

- **Data Source for AMI:**

```

data "aws_ami" "latest-amazon-linux-image" {
  most_recent = true
  owners     = ["amazon"]
  filter {
    name   = "name"
    values = ["Deep Learning Base AMI (Amazon Linux 2) Version 57.9"]
  }
}

output "aws-ami_id" {

```

```
value = data.aws_ami.latest-amazon-linux-image.id
}
```

Step 6: Create SSH Key Pair

- **Key Pair Resource:**

```
resource "aws_key_pair" "ssh-key" {
  key_name  = "server-key"
  public_key = file(var.public_key_location)
}
```

Step 7: Create the EC2 Instance

- **EC2 Instance Resource:**

```
resource "aws_instance" "myapp-server" {
  ami                = data.aws_ami.latest-amazon-linux-image.id
  instance_type      = var.instance_type
  subnet_id          = aws_subnet.myapp-subnet-1.id
  vpc_security_group_ids = [aws_default_security_group.default-sg.id]
  availability_zone   = var.avail_zone
  associate_public_ip_address = true
  key_name            = aws_key_pair.ssh-key.key_name

  user_data = file("entry-script.sh")
  user_data_replace_on_change = true

  connection {
    type      = "ssh"
    host      = self.public_ip
    user      = "ec2-user"
    private_key = file(var.private_key_location)
  }
}
```

```

}

provisioner "file" {
  source    = "entry-script.sh"
  destination = "/home/ec2-user/entry-script-on-ec2.sh"
}

provisioner "remote-exec" {
  script = "entry-script.sh"
}

provisioner "local-exec" {
  command = "echo ${self.public_ip} > output.txt"
}

tags = {
  Name = "${var.env_prefix}-server"
}

output "ec-public_ip" {
  value = aws_instance.myapp-server.public_ip
}

```

- **Entry Script (entry-script.sh):**

Create a file named **entry-script.sh** with the following content:

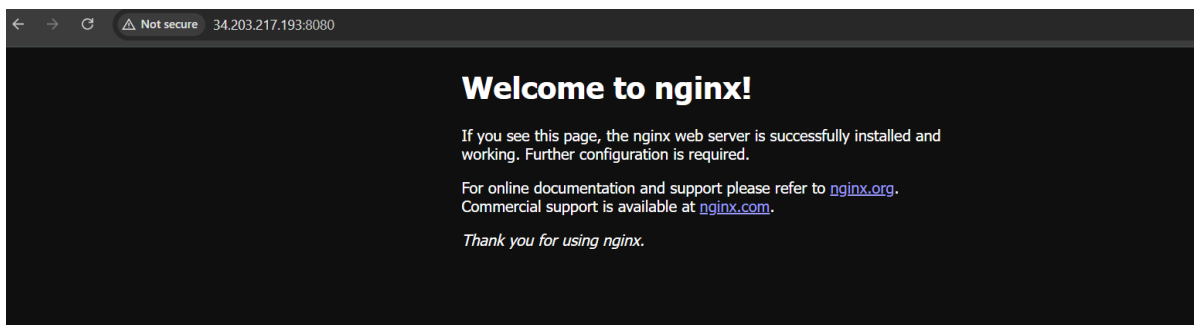
```

#!/bin/bash
sudo yum update -y
sudo yum install -y docker
sudo systemctl start docker
sudo usermod -aG docker ec2-user
docker run -p 8080:80 nginx

```

Step 8: Deploy the Infrastructure

1. **Preview the Terraform Plan:** `terraform plan`
2. **Apply the Configuration:** `terraform apply -auto-approve`
3. **SSH into the Instance:**
 - Since `~/.ssh/id_rsa` is the default location we can also do: `ssh ec2-user@<public ip>`
4. **Verify the Docker Container:** `docker ps`
 - Confirm that the Nginx container is running.
5. **Access the Application:**
 - Open your browser and navigate to: `http://<public-ip>:8080`



Step 9: Provisioners Branch (Feature: feature/provisioners)

In case you want to add additional provisioning steps (such as a custom script execution) you can create a separate feature branch called **feature/provisioners**. Provisioners are used as a last resort and are not recommended for primary configuration, but can be useful for post-deployment tasks.

1. **Create the Feature Branch:** `git checkout -b feature/provisioners`
2. **Add/Update Provisioner Blocks in Your Terraform Code:**
 - For example, you can add a remote-exec provisioner to execute additional configuration commands:

```
provisioner "remote-exec" {  
  inline = [  
    "echo 'Additional configuration or scripts can be executed here'",  
    "sudo docker ps"  
  ]  
}
```

Step 10: Clean Up

- **Destroy the Infrastructure:** `terraform destroy -auto-approve`
-