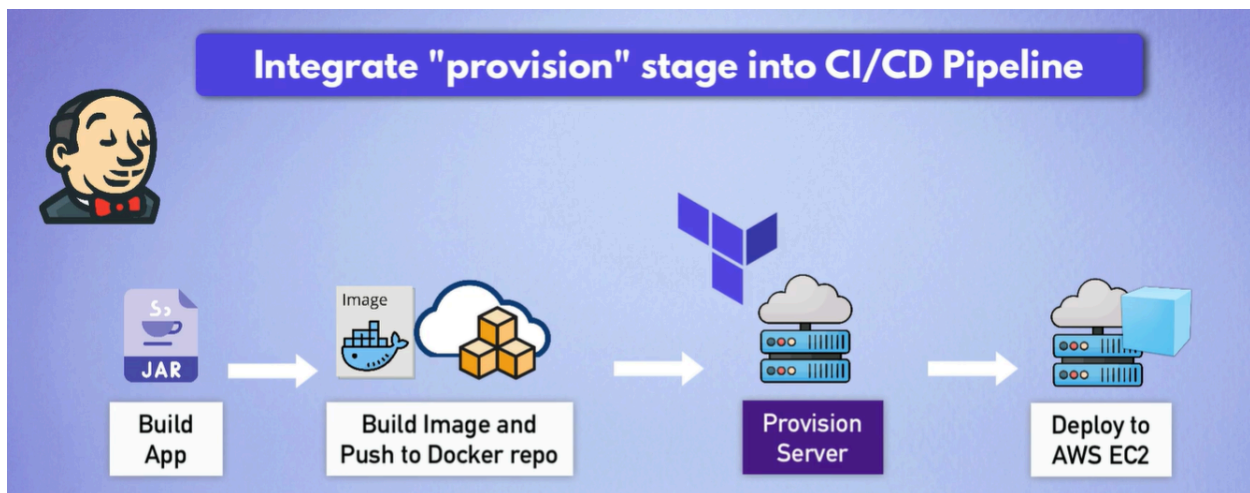


Demo Project: Complete CI/CD with Terraform

This guide outlines how to integrate a provisioning stage into a CI/CD pipeline to automate server provisioning and deployment using Jenkins and Terraform.

Stages

- **CI Step:** Build the artifact for a Java Maven application.
- **CI Step:** Build and push a Docker image.
- **CD Step:** Provision an EC2 instance using Terraform.
- **CD Step:** Deploy the application using Docker Compose.



Stages

Step 1: Create SSH key-pair

Step 2: Install Terraform in Jenkins container

Step 3: Add Terraform Configuration to Application Repository

Step 4: Add Provision Stage to Jenkinsfile

Step 5: Add Deploy Stage to Jenkinsfile

Step 6: Create `server-cmds.sh`

Step 7: Run CI/CD Pipeline

Step 8: Clean Up Resources

Troubleshooting

Issue: `wget: command not found`

Step 1: Create SSH key-pair

1. Navigate to the AWS EC2 Key Pair service.
2. Create a new key pair named `myapp-key-pair`.
3. Download the `.pem` file.
4. Add the key to Jenkins:
 - Open the Jenkins UI.
 - Go to **Manage Jenkins > Credentials**.
 - Select **Kind**: SSH Username with private key.
 - Username: `ec2-user`.
 - Private Key: Paste the content from `cat myapp-key-pair.pem`.

Dashboard > java-maven-app > Credentials > Folder > Global credentials (unrestricted) >

New credentials

Kind

SSH Username with private key

ID ?

server-ssh-key

Description ?

Username

ec2-user

☐ Treat username as secret ?

Step 2: Install Terraform in Jenkins container

1. SSH into the Jenkins host: `ssh root@<host-ip>`
2. Access the Jenkins container as root: `docker exec -it -u 0 <container-id> bash`
3. Check the Operating system being used: `cat /etc/os-release`

Example:

```
PRETTY_NAME="Debian GNU/Linux 12 (bookworm)"
NAME="Debian GNU/Linux"
VERSION_ID="12"
VERSION="12 (bookworm)"
VERSION_CODENAME=bookworm
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
```

4. Install Terraform

Download Terraform from <https://developer.hashicorp.com/terraform/install> based on your operating system.

Remember not to use `sudo` since we are login as root user.

- `wget -O - https://apt.releases.hashicorp.com/gpg | gpg --dearmor -o /usr/share/keyrings/hashicorp-archive-keyring.gpg`
- `echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/hashicorp-archive-keyring.gpg] https://apt.releases.hashicorp.com $(grep VERSION_CODENAME /etc/os-release | cut -d= -f2) main" | tee /etc/apt/sources.list.d/hashicorp.list`
- `apt update && apt install terraform`

5. Verify installation: `terraform -v`

Step 3: Add Terraform Configuration to Application Repository

1. In your application's Git repository, create a `terraform/` directory.
 2. Add `main.tf` with the Terraform configuration.
 3. Create a `variables.tf` file in the same directory to define variables.
-

Step 4: Add Provision Stage to Jenkinsfile

Add a new stage to the Jenkinsfile:

```
stage("provision server") {  
  environment {  
    AWS_ACCESS_KEY_ID = credentials('jenkins_aws_access_key_id')  
    AWS_SECRET_ACCESS_KEY = credentials('jenkins-aws_secret_access_key')  
    TF_VAR_env_prefix = 'test'// Override variable  
  }  
  steps {  
    script {  
      dir('terraform') {  
        sh "terraform init"  
        sh "terraform apply --auto-approve"  
        EC2_PUBLIC_IP = sh(  
          script: "terraform output ec2-public_ip",  
          returnStdout: true  
        ).trim()  
      }  
    }  
  }  
}
```

Step 5: Add Deploy Stage to Jenkinsfile

Ensure the deployment stage dynamically retrieves the EC2 public IP. We don't want to hard-code the EC2 public IP since Terraform outputs it into an environment variable, `EC2_PUBLIC_IP`. Access it as shown below:

```
stage("deploy") {
  environment {
    DOCKER_CREDS = credentials('docker-hub-repo')
  }
  steps {
    script {
      echo "waiting for EC2 server to initialize"
      sleep(time: 90, unit: "SECONDS")

      echo 'deploying docker image to EC2...'
      echo "${EC2_PUBLIC_IP}"

      def shellCmd = "bash ./server-cmds.sh ${IMAGE_NAME} ${DOCKER_CREDS}"
      def ec2Instance = "ec2-user@${EC2_PUBLIC_IP}"

      sshagent(['server-ssh-key']) {
        sh "scp -o StrictHostKeyChecking=no server-cmds.sh ${ec2Instance}:/home/ec2-user/"
        sh "scp -o StrictHostKeyChecking=no docker-compose.yaml ${ec2Instance}:/home/ec2-user/"
        sh "ssh -o StrictHostKeyChecking=no ${ec2Instance} ${shellCmd}"
      }
    }
  }
}
```

Step 6: Create `server-cmds.sh`

This script will log in to Docker and start the application. The script takes three arguments:

`IMAGE_NAME` , `DOCKER_CREDS_USR` , and `DOCKER_CREDS_PSW`

```
#!/usr/bin/env bash
```

```
export IMAGE=$1
export DOCKER_USER=$2
export DOCKER_PWD=$3
echo $DOCKER_PWD | docker login -u $DOCKER_USER --password-stdin
docker-compose -f docker-compose.yaml up --detach
echo "success"
```

In Jenkinsfile:

```
...
def shellCmd = "bash ./server-cmds.sh ${IMAGE_NAME} ${DOCKER_CREDS_USR}
...
```

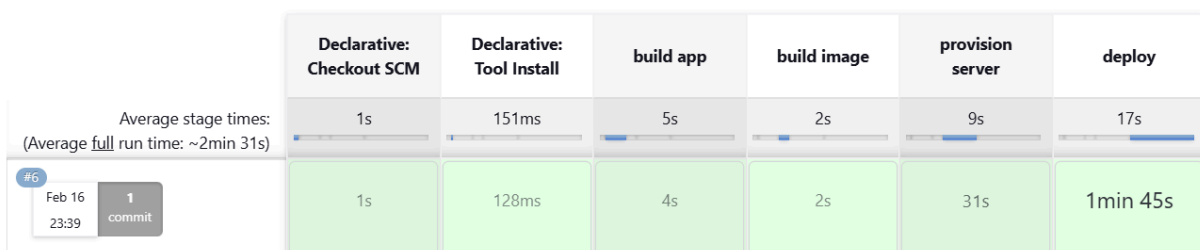
Step 7: Run CI/CD Pipeline

1. Commit and push changes to the Git repository.
2. Check the Jenkins pipeline and ensure all stages complete successfully.

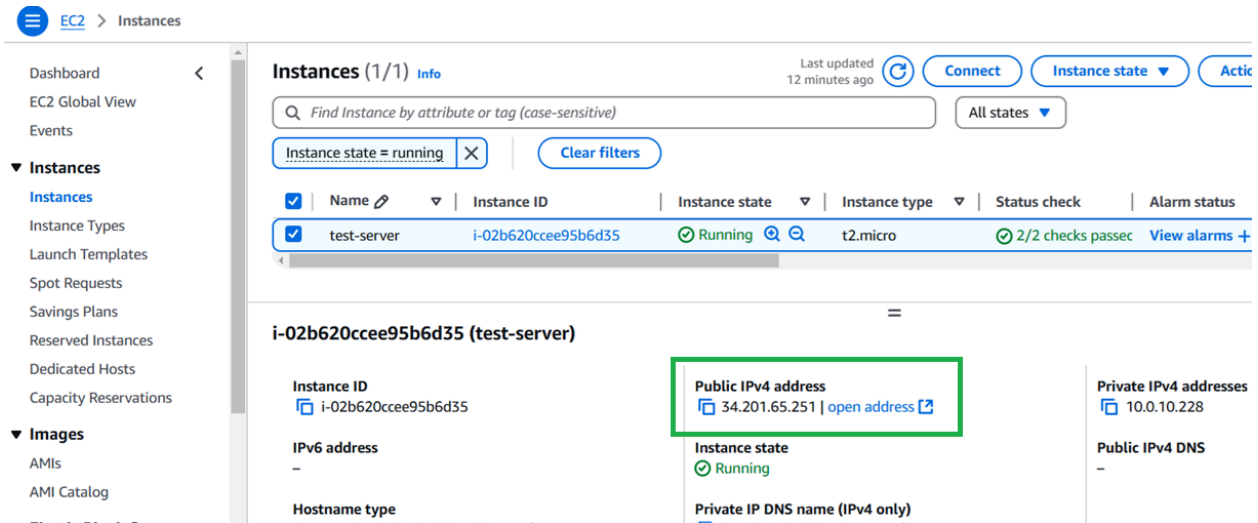
✓ jenkinsfile-sshagent

Full project name: java-maven-app/jenkinsfile-sshagent

Stage View



3. Retrieve the public IP of the provisioned EC2 instance:



The screenshot shows the AWS Management Console for EC2 Instances. The left sidebar contains navigation links for Dashboard, EC2 Global View, Events, and a list of instance-related services. The main content area displays a table of instances. One instance, 'test-server' (ID: i-02b620ccee95b6d35), is shown in a 'Running' state. Below the table, the details for this instance are expanded, showing its public IPv4 address as 34.201.65.251, which is highlighted with a green box. Other details include the instance type (t2.micro), status check (2/2 checks passed), and private IP address (10.0.10.228).

4. SSH into the instance:

```
chmod 400 myapp-key-pair.pem
ssh -i /path/to/myapp-key-pair.pem ec2-user@<ec2-public-ip>
```

5. Verify the running containers: `docker ps`

Example output:

```
[ec2-user@ip-10-0-10-228 ~]$ docker ps
```

| CONTAINER ID | IMAGE | COMMAND | CREATED |
|--------------|--|--------------------------|---------|
| 844f5b915927 | postgres:15 | "docker-entrypoint.s..." | 12 mi |
| e70fbe389cbf | eduardobautistamaciell/demo-app:java-maven-2.0 | "/bin/sh -c 'j | |

Note: Remove `sudo` if you are the root user.
